

Checking membership

CLEANING DATA IN R



Maggie Matsui

Content Developer @ DataCamp

Categorical data

- Categorical variables have a fixed and known set of possible values

Data	Example values
Marriage status	unmarried , married
Household income category	0-20K , 20-40K , ...
T-shirt size	S , M , L , XL

Factors

- In a `factor`, each category is stored as a number and has a corresponding label

Data	Labels	Numeric representation
Marriage status	<code>unmarried</code> , <code>married</code>	<code>1</code> , <code>2</code>
Household income category	<code>0-20K</code> , <code>20-40K</code> , ...	<code>1</code> , <code>2</code> , ...
T-shirt size	<code>S</code> , <code>M</code> , <code>L</code> , <code>XL</code>	<code>1</code> , <code>2</code> , <code>3</code> , <code>4</code>

Factor levels

```
tshirt_size
```

```
L  XL XL L  M  M  M  L  XL L  S  M  M  S  S  M  XL S  L  S ...  
Levels: S M L XL
```

```
levels(tshirt_size)
```

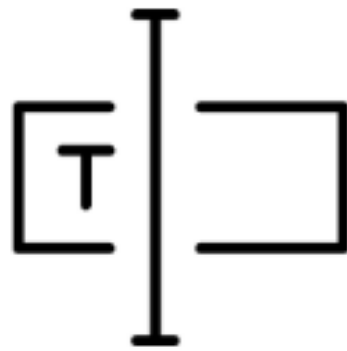
```
"S"  "M"  "L"  "XL"
```

Values that don't belong

- `factor`s cannot have values that fall outside of the predefined ones

Data	Levels	Not allowed
Marriage status	<code>unmarried</code> , <code>married</code>	<code>divorced</code>
Household income category	<code>0-20K</code> , <code>20-40K</code> , ...	<code>10-30K</code>
T-shirt size	<code>S</code> , <code>M</code> , <code>L</code> , <code>XL</code>	<code>S/M</code>

How do we end up with these values?



Free text

Or



Dropdowns

Data Entry Errors



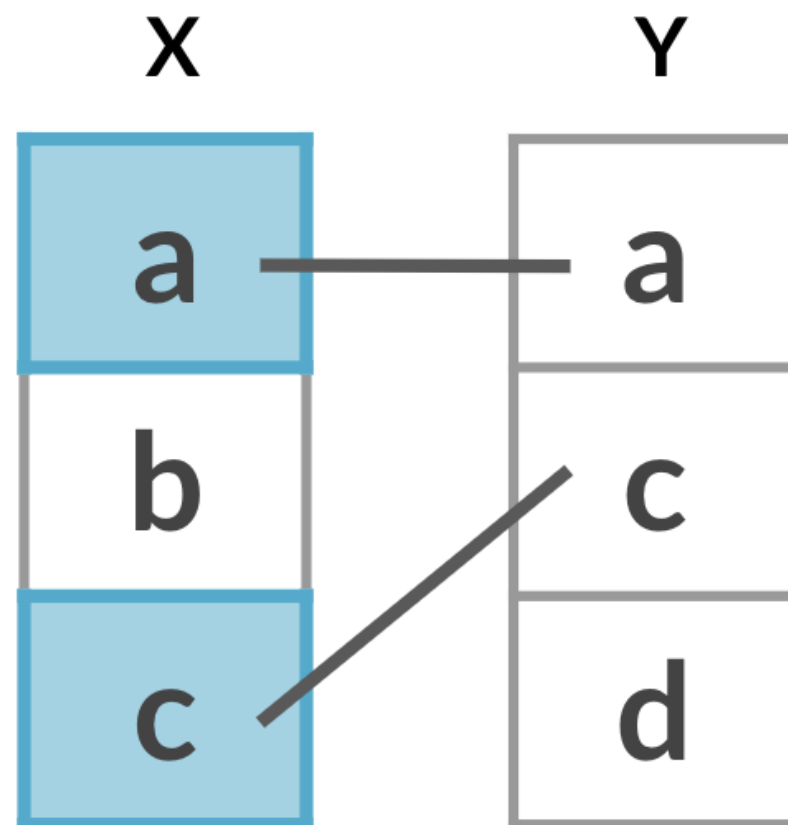
Parsing Errors

Filtering joins: a quick review

- Keeps or removes observations from the first table without adding columns

Semi-join

*What observations of X
are **also** in Y?*

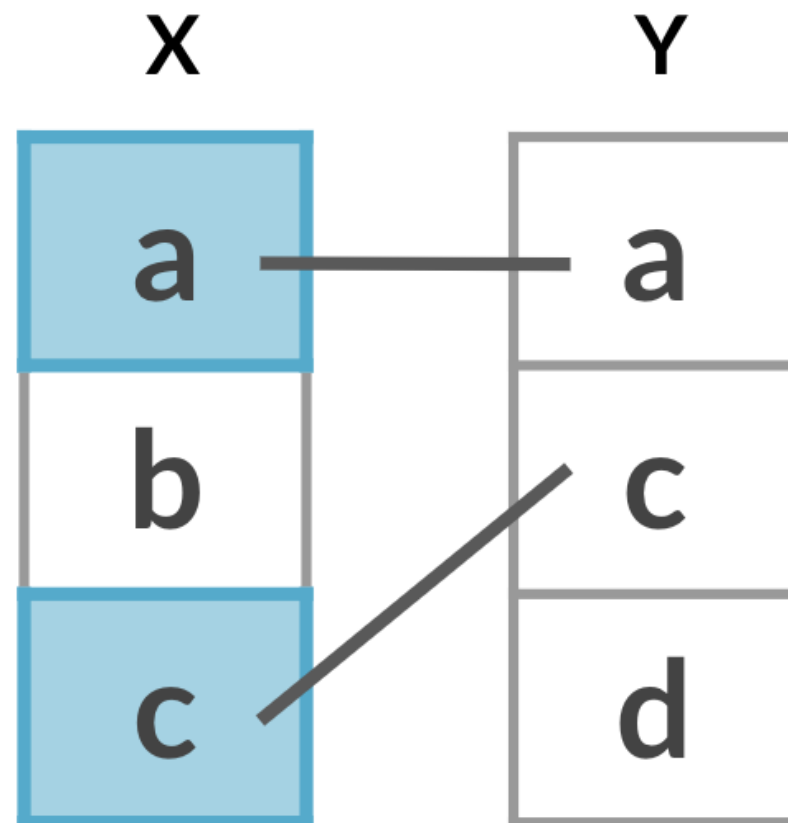


Filtering joins: a quick review

- Keeps or removes observations from the first table without adding columns

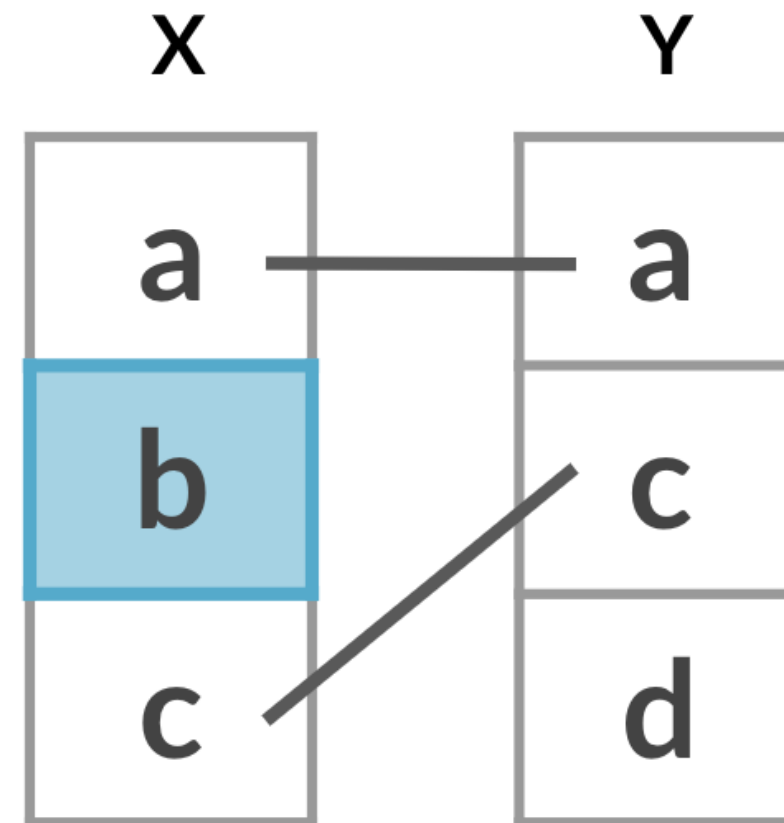
Semi-join

*What observations of X
are **also** in Y?*



Anti-join

*What observations of X
are **not** in Y?*



Blood type example

study_data

	name	birthday	blood_type
1	Beth	2019-10-20	B-
2	Ignatius	2020-07-08	A-
3	Paul	2019-08-12	O+
4	Helen	2019-03-17	O-
5	Jennifer	2019-12-17	Z+
6	Kennedy	2020-04-27	A+
7	Keith	2019-04-19	AB+

blood_types

	blood_type
1	O-
2	O+
3	A-
4	A+
5	B+
6	B-
7	AB+
8	AB-

Blood type example

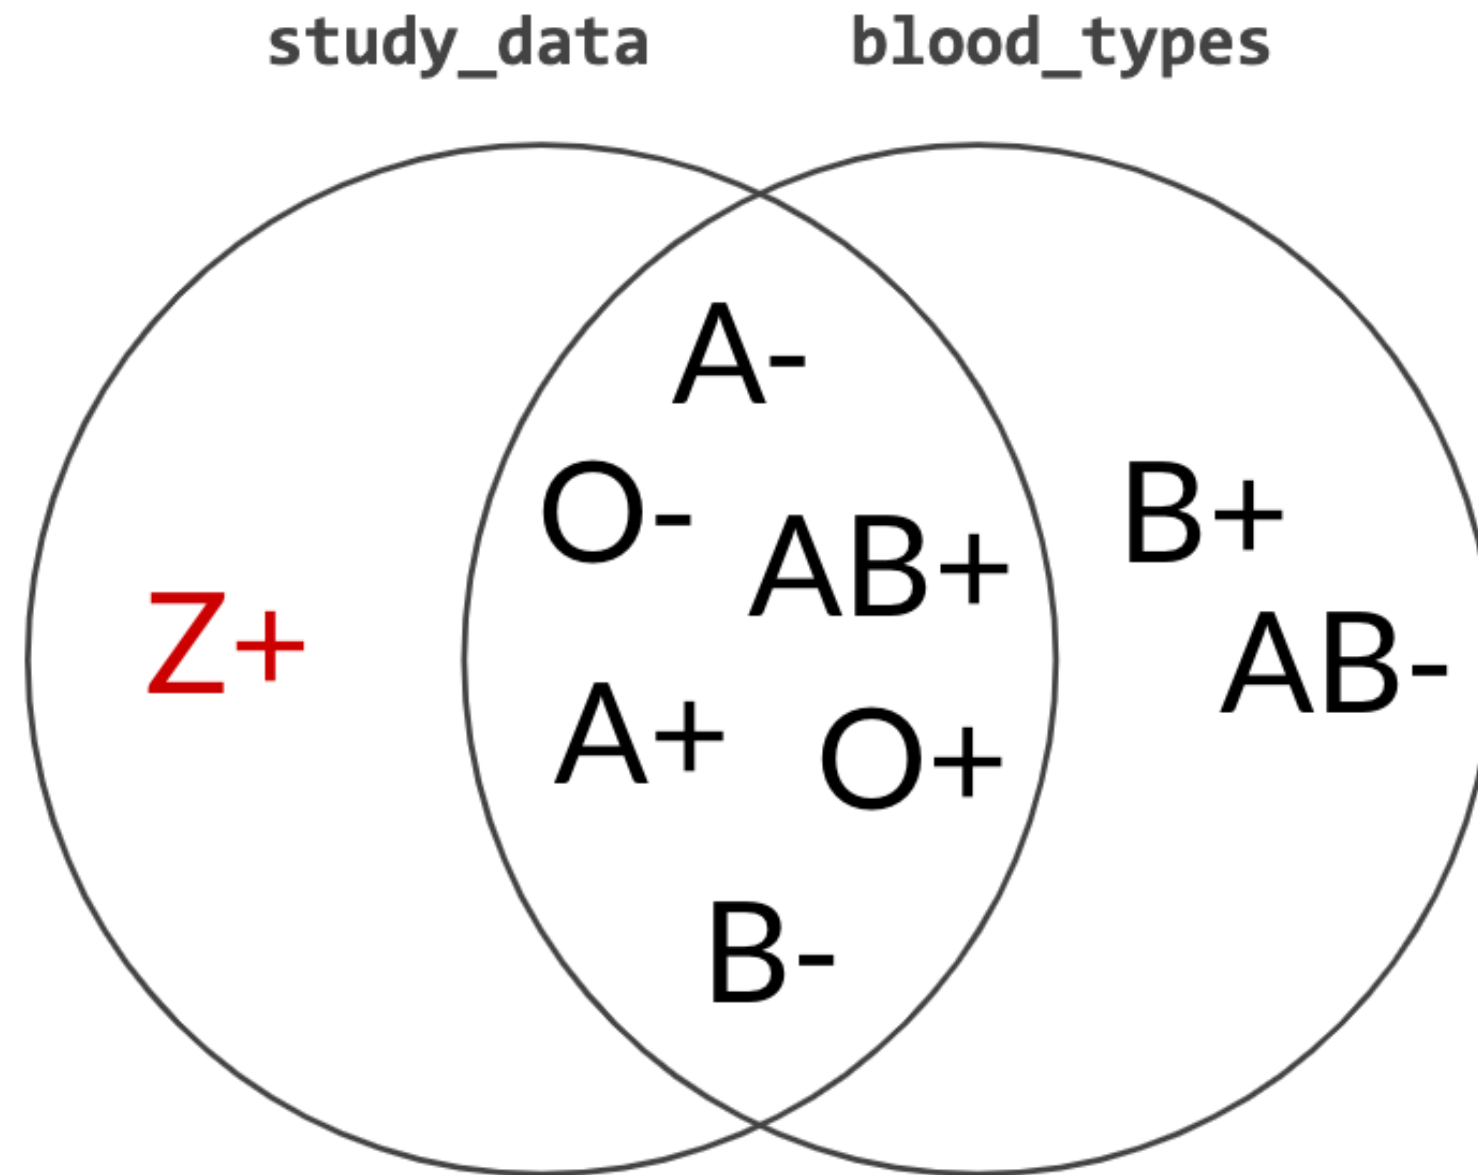
study_data

	name	birthday	blood_type
1	Beth	2019-10-20	B-
2	Ignatius	2020-07-08	A-
3	Paul	2019-08-12	O+
4	Helen	2019-03-17	O-
5	Jennifer	2019-12-17	Z+ <--
6	Kennedy	2020-04-27	A+
7	Keith	2019-04-19	AB+

blood_types

	blood_type
1	O-
2	O+
3	A-
4	A+
5	B+
6	B-
7	AB+
8	AB-

Finding non-members

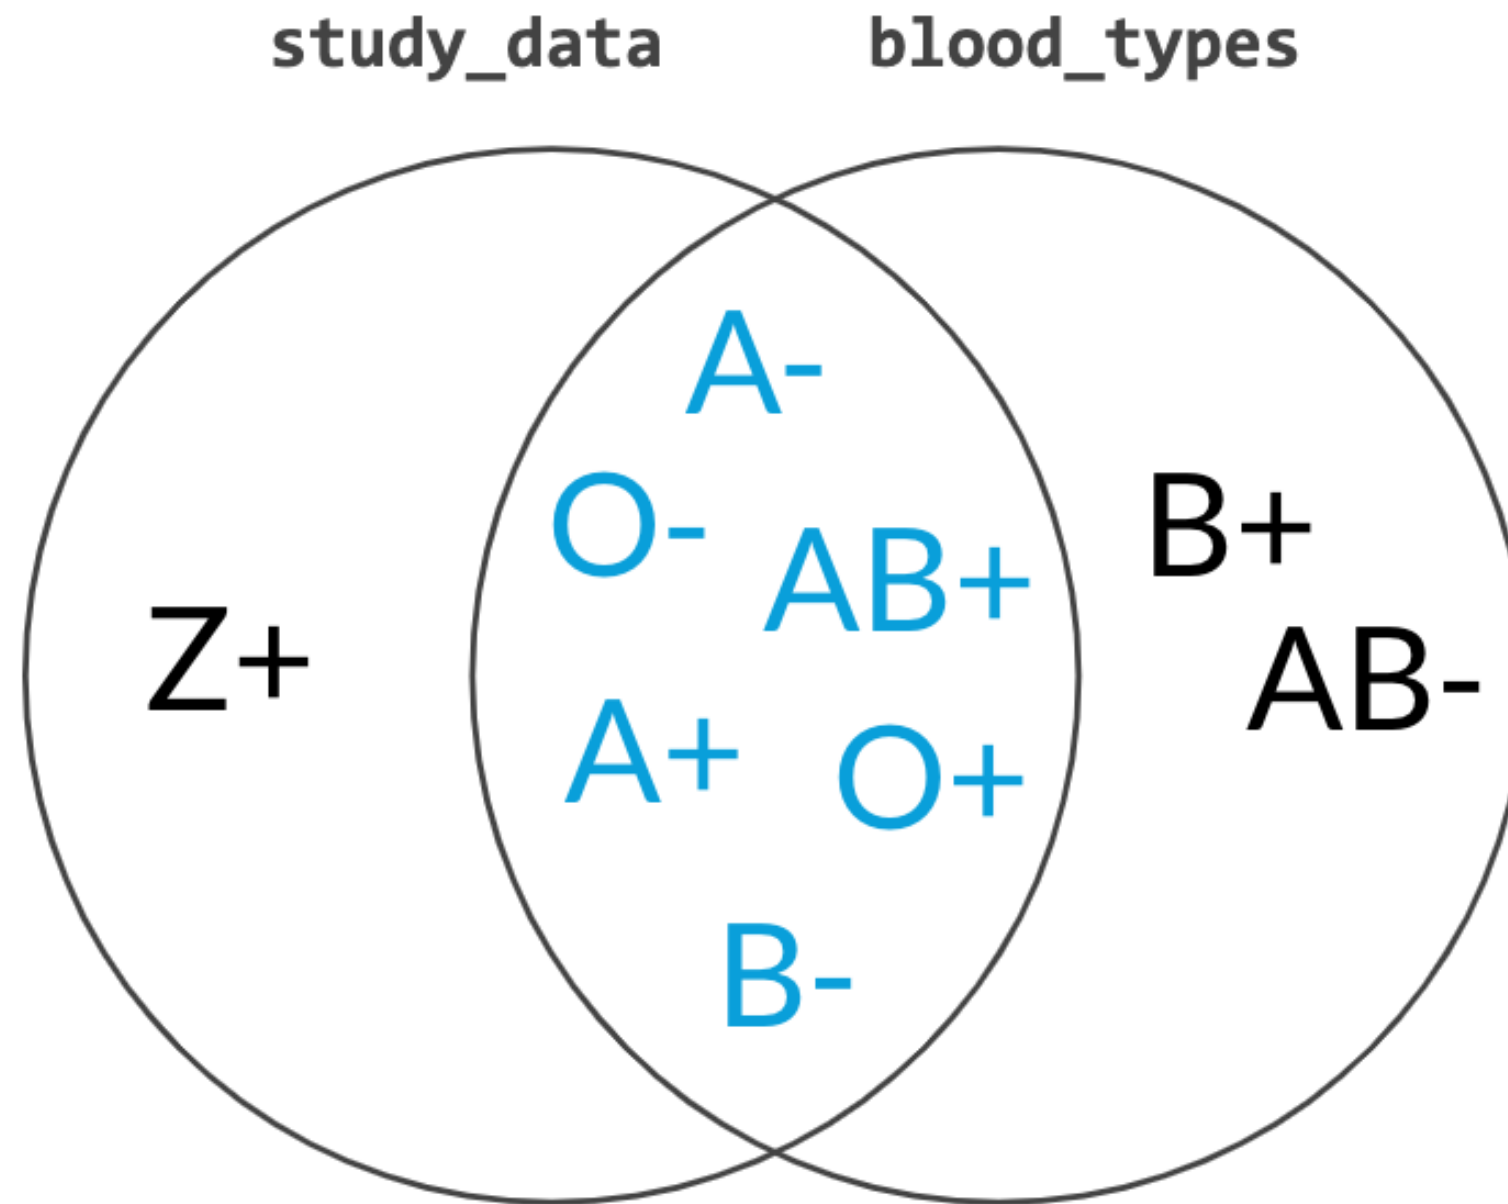


Anti-join

```
study_data %>%  
  anti_join(blood_types, by = "blood_type")
```

```
   name    birthday blood_type  
1 Jennifer 2019-12-17        Z+
```

Removing non-members



Semi-join

```
study_data %>%  
  semi_join(blood_types, by = "blood_type")
```

	name	birthday	blood_type
1	Beth	2019-10-20	B-
2	Ignatius	2020-07-08	A-
3	Paul	2019-08-12	O+
4	Helen	2019-03-17	O-
5	Kennedy	2020-04-27	A+
6	Keith	2019-04-19	AB+

Let's practice!
CLEANING DATA IN R

Categorical data problems

CLEANING DATA IN R

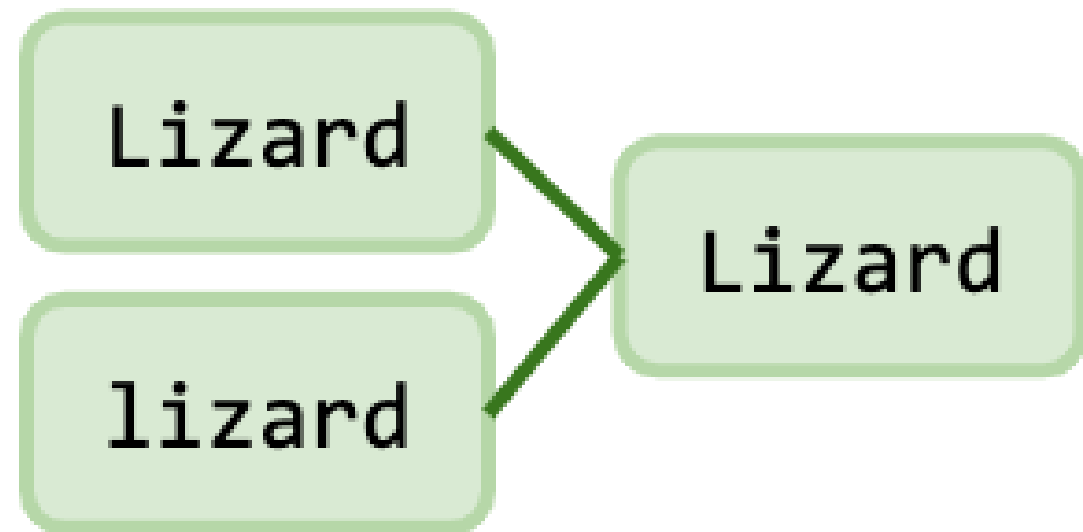


Maggie Matsui

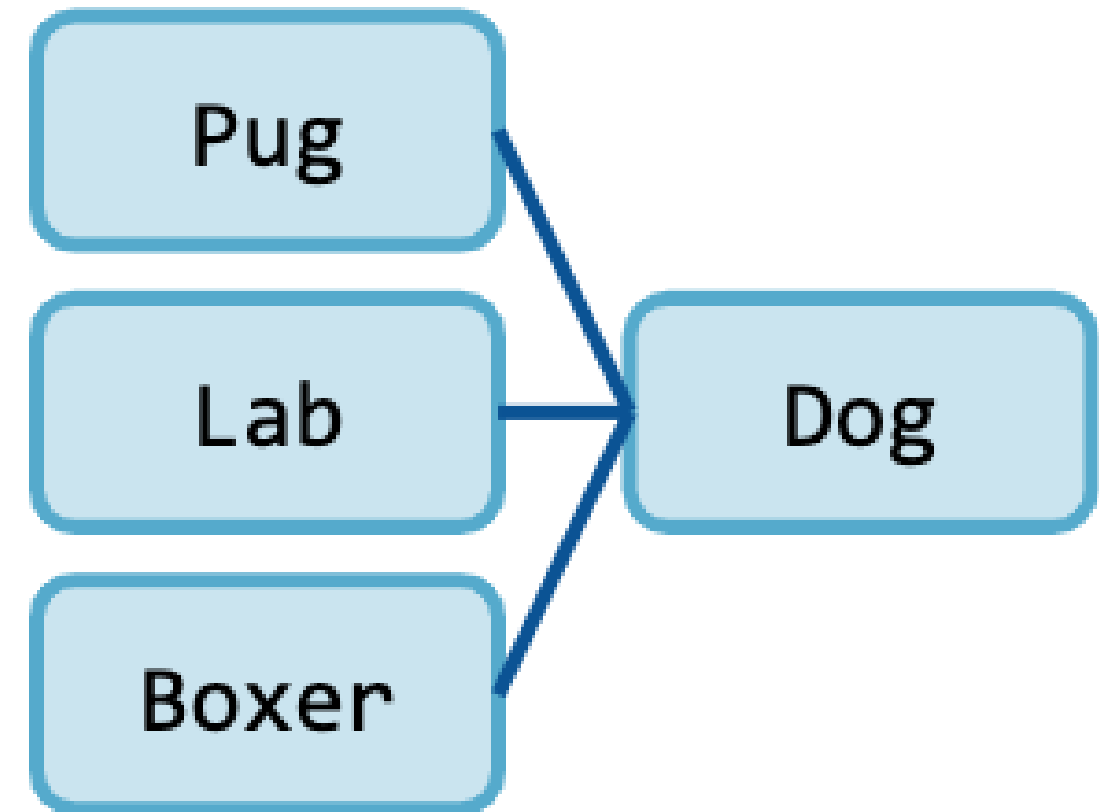
Content Developer @ DataCamp

Categorical data problems

Inconsistency within a category



Too many categories



Example: animal classification

```
animals
```

```
# A tibble: 68 x 9
  animal_name hair  eggs  fins  legs tail  type
  <chr>      <fct> <fct> <fct> <int> <fct> <fct>
1 mole      1      0      0      4 1     mammal
2 chicken   0      1      0      2 1     bird
3 capybara  1      0      0      2 1     Mammal
4 tuna      0      1      1      0 1     fish
5 ostrich   0      1      0      2 1     bird
# ... with 63 more rows
```

Checking categories

```
animals %>%  
  count(type)
```

- "mammal"
- " mammal "
- "MAMMAL"
- "Mammal "

	type	n
1	" mammal "	1
2	"amphibian"	2
3	"bird"	20
4	"bug"	1
5	"fish"	2
6	"invertebrate"	1
7	"mammal"	38
8	"MAMMAL"	1
9	"Mammal "	1
10	"reptile"	1

Case inconsistency

```
library(stringr)
animals %>%
  mutate(type_lower = str_to_lower(type))
```

	animal_name	hair	eggs	fins	legs	tail	type	type_lower
	<fct>	<int>	<int>	<int>	<int>	<int>	<fct>	<chr>
1	mole	1	0	0	4	1	"mammal"	"mammal"
2	chicken	0	1	0	2	1	"bird"	"bird"
3	capybara	1	0	0	2	1	" Mammal"	" mammal"
4	tuna	0	1	1	0	1	"fish"	"fish"
5	ostrich	0	1	0	2	1	"bird"	"bird"

Case inconsistency

```
animals %>%  
  mutate(type_lower = str_to_lower(type)) %>%  
  count(type_lower)
```

	type_lower	n		type_lower	n
	<chr>	<int>		<chr>	<int>
1	"mammal "	1	6	"invertebrate"	1
2	"amphibian"	2	7	"mammal"	39
3	"bird"	20	8	"mammal "	1
4	"bug"	1	9	"reptile"	1
5	"fish"	2			

"MAMMAL" → "mammal"

Case inconsistency

```
animals %>%  
  mutate(type_upper = str_to_upper(type)) %>%  
  count(type_upper)
```

	type_upper	n		type_upper	n
	<chr>	<int>		<chr>	<int>
1	" MAMMAL "	1	6	"INVERTEBRATE"	1
2	"AMPHIBIAN"	2	7	"MAMMAL"	39
3	"BIRD"	20	8	"MAMMAL "	1
4	"BUG"	1	9	"REPTILE"	1
5	"FISH"	2			

Whitespace inconsistency

```
animals %>%  
  mutate(type_trimmed = str_trim(type_lower))
```

	animal_name	hair	eggs	fins	legs	tail	type_lower	type_trimmed
	<fct>	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>
1	mole	1	0	0	4	1	"mammal"	mammal
2	chicken	0	1	0	2	1	"bird"	bird
3	capibara	1	0	0	2	1	" mammal"	mammal
4	tuna	0	1	1	0	1	"fish"	fish
5	ostrich	0	1	0	2	1	"bird"	bird

Whitespace inconsistency

```
animals %>%  
  mutate(type_trimmed = str_trim(type_lower)) %>%  
  count(type_trimmed)
```

	type_trimmed	n		type_trimmed	n
	<chr>	<int>		<chr>	<int>
1	amphibian	2	6	mammal	41
2	bird	20	7	reptile	1
3	bug	1			
4	fish	2			
5	invertebrate	1			

Too many categories

```
animals %>%  
  count(type_trimmed, sort = TRUE)
```

	type_trimmed	n
1	mammal	41
2	bird	20
3	amphibian	2
4	fish	2
5	bug	1
6	invertebrate	1
7	reptile	1

Collapsing categories

```
other_categories = c("amphibian", "fish", "bug", "invertebrate", "reptile")
```

```
library(forcats)
```

```
animals %>%
```

```
  mutate(type_collapsed = fct_collapse(type_trimmed, other = other_categories))
```

	animal_name	hair	eggs	fins	legs	tail	type_trimmed	type_collapsed
	<fct>	<int>	<int>	<int>	<int>	<int>	<chr>	<chr>
1	mole	1	0	0	4	1	mammal	mammal
2	chicken	0	1	0	2	1	bird	bird
3	capibara	1	0	0	2	1	mammal	mammal
4	tuna	0	1	1	0	1	fish	other
5	ostrich	0	1	0	2	1	bird	bird

Collapsing categories

```
animals %>%  
  count(type_collapsed)
```

```
type_collapsed      n  
<fct>             <int>  
1 other              7  
2 bird             20  
3 mammal            41
```

```
animals %>%  
  group_by(type_collapsed) %>%  
  summarize(avg_legs = mean(legs))
```

```
type_collapsed avg_legs  
<fct>          <dbl>  
1 other         3.71  
2 bird           2  
3 mammal        3.37
```

Let's practice!
CLEANING DATA IN R

Cleaning text data

CLEANING DATA IN R



Maggie Matsui

Content Developer @ DataCamp

What is text data?

Type of data	Example values
Names	"Veronica Hopkins" , "Josiah" , ...
Phone numbers	"6171679912" , "(868) 949-4489" , ...
Emails	"vhopkins@datacamp.com" , "josiah@josiah.co.uk" , ...
Passwords	"JZY46TVG8SM" , "iamjosiah21" , ...
Comments/Reviews	"great service!" , "This product broke after 2 days" , ...

Unstructured data problems

- Formatting inconsistency
 - "6171679912" vs. "(868) 949-4489"
 - "9239 5849 3712 0039" vs. "4490459957881031"
- Information inconsistency
 - +1 617-167-9912 vs. 617-167-9912
 - "Veronica Hopkins" vs. "Josiah"
- Invalid data
 - Phone number "0492" is too short
 - Zip code "19888" doesn't exist

Customer data

```
customers
```

```
# A tibble: 99 x 3
  name                company                credit_card
  <chr>                <chr>                <chr>
1 Galena              In Magna Associates  5171 5854 8986 1916
2 MacKenzie           Iaculis Ltd         5128-5078-8008-5824
3 Megan Acosta        Semper LLC           5502 4529 0732 1744
4 Phoebe Delacruz    Sit Amet Nulla Limited 5419-7308-7424-0944
5 Jessica             Pellentesque Sed Ltd  5419 2949 5508 9530
# ... with 95 more rows
```


Detecting hyphenated credit card numbers

```
str_detect(customers$credit_card, "-")
```

```
FALSE TRUE FALSE TRUE FALSE TRUE TRUE FALSE FALSE TRUE TRUE TRUE FALSE ...
```

```
customers %>%  
  filter(str_detect(credit_card, "-"))
```

	name	company	credit_card
1	MacKenzie	Iaculis Ltd	5128-5078-8008-5824
2	Phoebe Delacruz	Sit Amet Nulla Limited	5419-7308-7424-0944
3	Abel	Lorem PC	5211-6023-0805-0217
...			

Replacing hyphens

```
customers %>%  
  mutate(credit_card_spaces = str_replace_all(credit_card, "-", " "))
```

	name	company	credit_card_spaces
1	Galena	In Magna Associates	5171 5854 8986 1916
2	MacKenzie	Iaculis Ltd	5128 5078 8008 5824
3	Megan Acosta	Semper LLC	5502 4529 0732 1744
4	Phoebe Delacruz	Sit Amet Nulla Limited	5419 7308 7424 0944
5	Jessica	Pellentesque Sed Ltd	5419 2949 5508 9530
...			

Removing hyphens and spaces

```
credit_card_clean <- customers$credit_card %>%  
  str_remove_all("-") %>%  
  str_remove_all(" ")  
customers %>%  
  mutate(credit_card = credit_card_clean)
```

	name	company	credit_card
1	Galena	In Magna Associates	5171585489861916
2	MacKenzie	Iaculis Ltd	5128507880085824
3	Megan Acosta	Semper LLC	5502452907321744
...			

Finding invalid credit cards

```
str_length(customers$credit_card)
```

```
16 16 16 16 16 16 16 16 16 16 16 16 12 16 16 16 16 16 16 16 16 16 16 16 ...
```

```
customers %>%  
  filter(str_length(credit_card) != 16)
```

	name	company	credit_card
1	Jerry Russell	Sed Eu Company	516294099537
2	Ivor Christian	Ut Tincidunt Incorporated	544571330015
3	Francesca Drake	Etiam Consulting	517394144089

Removing invalid credit cards

```
customers %>%  
  filter(str_length(credit_card) == 16)
```

	name	company	credit_card
1	Galena	In Magna Associates	5171585489861916
2	MacKenzie	Iaculis Ltd	5128507880085824
3	Megan Acosta	Semper LLC	5502452907321744
4	Phoebe Delacruz	Sit Amet Nulla Limited	5419730874240944
5	Jessica	Pellentesque Sed Ltd	5419294955089530
...			

More complex text problems

- A *regular expression* is a sequence of characters that allows for robust searching within a string.
- Certain characters are treated differently in a regular expression:
 - `(`, `)`, `[`, `]`, `$`, `.`, `+`, `*`, and others
- `stringr` functions use regular expressions
- Searching for these characters requires using `fixed()` :
 - `str_detect(column, fixed("$"))`

Learn more in [String Manipulation with stringr in R](#) & [Intermediate Regular Expressions in R](#)

Let's practice!
CLEANING DATA IN R